

Engineered Intelligence workflow and audit boundaries

What EI does: train/CV candidate discovery -> sparse fusion -> calibration -> locked holdout score



Group-locked split

train/CV rows only; holdout scored once

Claim boundary

synthetic public-proxy evidence, not measured truth

Feature denylist

labels, group IDs, split keys, and audit headers blocked

Code boundary

workflow reviewable; generated arrays stay out of Git

Raw component handles, selection_lock.json, and component CSVs remain evidence artifacts; reader-facing labels summarize their role.